

Final Project  
Case Study in Finance - House Rooms Classification

Rachapudi, Vinod



CSCI E-89 Deep Learning, Fall 2023  
**Harvard University Extension School**  
Prof. Zoran B. Djordjević

# Introduction

In FinTech Mortgage business investment (Mortgage Portfolio) decisions are made based on a house and its rooms.

Therefore, having the ability to classify the rooms of a house to identify the type of room such as Dining Room vs Bed Room vs et al enables these investment decisions.



The code in this project helps with binary classification of Dining Room vs Bed Room by using

- **Convolutional Neural Networks in Part 1**
- **Pre-trained Convolutional Neural Network VGG16 in Part 2 and**
- **Pre-trained Convolutional Neural Networks - Xception in Part 3**

# Downloading the Data

The Data is downloaded from Kaggle <https://www.kaggle.com/datasets/robinreni/house-rooms-image-dataset/data>

This Kaggle Data set has 1248 bedroom pictures and 1158 dining room pictures



# Splitting of Data into Training, Validation & Test DataSets

This Kaggle Data set has 1248 bedroom pictures and 1158 dining room pictures

Total training bedroom images: 500

Total training diningroom images: 500

Total validation bedroom images: 300

Total validation diningroom images: 300

Total test bedroom images: 448

Total test diningroom images: 358

# Part 1 - Convolutional Neural Networks



# Conv Neural Network

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513

=====  
Total params: 3453121 (13.17 MB)  
Trainable params: 3453121 (13.17 MB)  
Non-trainable params: 0 (0.00 Byte)

# Validation Accuracy is at 82.5%



Epoch 30/30

20/20 [=====] - 3s 136ms/step - loss: 0.1534 - acc: 0.9525  
- val\_loss: 0.6058 - val\_acc: 0.8250



# Added the L1 Regularizer (0.0001) on the Dense Layer

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3211776
dense_3 (Dense)	(None, 1)	513
Total params: 3453121 (13.17 MB)		
Trainable params: 3453121 (13.17 MB)		
Non-trainable params: 0 (0.00 Byte)		



# Regularization does not help 82.5% vs 82.2%



Epoch 15/15



20/20 [=====] - 3s 132ms/step - loss: 0.3307 - acc: 0.9000  
- val\_loss: 0.4830 - val\_acc: 0.8225

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 150, 150, 3)]	0	None
rescaling (Rescaling)	(None, 150, 150, 3)	0	['input_1[0][0]']
conv2d_8 (Conv2D)	(None, 146, 146, 32)	2400	['rescaling[0][0]']
batch_normalization (Batch Normalization)	(None, 146, 146, 32)	128	['conv2d_8[0][0]']
activation (Activation)	(None, 146, 146, 32)	0	['batch_normalization[0][0]']
separable_conv2d (Separable Conv2D)	(None, 146, 146, 32)	1312	['activation[0][0]']
batch_normalization_1 (Batch Normalization)	(None, 146, 146, 32)	128	['separable_conv2d[0][0]']
activation_1 (Activation)	(None, 146, 146, 32)	0	['batch_normalization_1[0][0]']
separable_conv2d_1 (Separable Conv2D)	(None, 146, 146, 32)	1312	['activation_1[0][0]']
max_pooling2d_8 (MaxPooling2D)	(None, 73, 73, 32)	0	['separable_conv2d_1[0][0]']
batch_normalization_2 (Batch Normalization)	(None, 73, 73, 32)	128	['max_pooling2d_8[0][0]']
activation_2 (Activation)	(None, 73, 73, 32)	0	['batch_normalization_2[0][0]']
separable_conv2d_2 (Separable Conv2D)	(None, 73, 73, 64)	2336	['activation_2[0][0]']
batch_normalization_3 (Batch Normalization)	(None, 73, 73, 64)	256	['separable_conv2d_2[0][0]']
activation_3 (Activation)	(None, 73, 73, 64)	0	['batch_normalization_3[0][0]']
separable_conv2d_3 (Separable Conv2D)	(None, 73, 73, 64)	4672	['activation_3[0][0]']
max_pooling2d_9 (MaxPooling2D)	(None, 37, 37, 64)	0	['separable_conv2d_3[0][0]']
batch_normalization_4 (Batch Normalization)	(None, 37, 37, 64)	256	['max_pooling2d_9[0][0]']
activation_4 (Activation)	(None, 37, 37, 64)	0	['batch_normalization_4[0][0]']
separable_conv2d_4 (Separable Conv2D)	(None, 37, 37, 128)	8768	['activation_4[0][0]']
batch_normalization_5 (Batch Normalization)	(None, 37, 37, 128)	512	['separable_conv2d_4[0][0]']
activation_5 (Activation)	(None, 37, 37, 128)	0	['batch_normalization_5[0][0]']
separable_conv2d_5 (Separable Conv2D)	(None, 37, 37, 128)	17536	['activation_5[0][0]']
max_pooling2d_10 (MaxPooling2D)	(None, 19, 19, 128)	0	['separable_conv2d_5[0][0]']
batch_normalization_6 (Batch Normalization)	(None, 19, 19, 128)	512	['max_pooling2d_10[0][0]']
activation_6 (Activation)	(None, 19, 19, 128)	0	['batch_normalization_6[0][0]']
separable_conv2d_6 (Separable Conv2D)	(None, 19, 19, 256)	33920	['activation_6[0][0]']
batch_normalization_7 (Batch Normalization)	(None, 19, 19, 256)	1024	['separable_conv2d_6[0][0]']
activation_7 (Activation)	(None, 19, 19, 256)	0	['batch_normalization_7[0][0]']
separable_conv2d_7 (Separable Conv2D)	(None, 19, 19, 256)	67840	['activation_7[0][0]']

# Replacing with Separable Conv2D

The validation accuracy went up to 86.25% with the SeparableConv2D replacement from the original (only Conv2D withOUT regularization) 82.5%



Epoch 60/60

20/20 [=====] - 3s 131ms/step - loss: 0.0567 - acc: 0.9825  
- val\_loss: 1.0717 - val\_acc: 0.8625

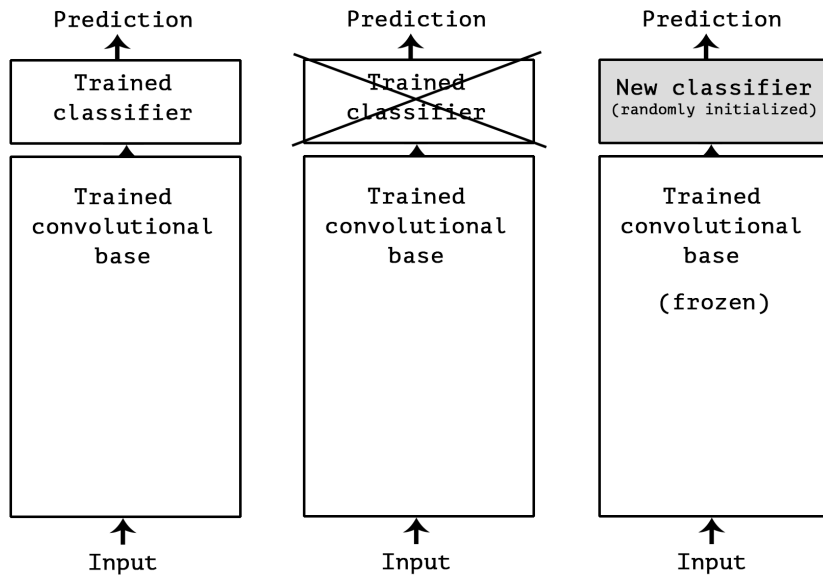
**This shows that Separable CONV2D leads to ~4% increase in accuracy which is a big deal. There is still some overfitting.**

**The of Trainable params decreased to 674,817 from the original**

**Trainable params: 3,453,121 - which is good as with fewer parameters we are getting a 4% accuracy lift.**

## Part 2 - Pre-trained Convolutional Neural Networks - VGG16

# Pre-trained Convolutional Neural Network VGG16



Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0

Total params: 14714688 (56.13 MB)  
 Trainable params: 14714688 (56.13 MB)  
 Non-trainable params: 0 (0.00 Byte)

## Densely-connected classifier



```
from keras import models
from keras import layers
from keras import optimizers

model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(train_features, train_labels,
                    epochs=30,
                    batch_size=20,
                    validation_data=(validation_features, validation_labels))
```

# Validation Accuracy 92.67%



Epoch 30/30

50/50 [=====] - 0s 8ms/step - loss: 0.0191 - acc: 0.9950 -  
val\_loss: 0.4455 - val\_acc: 0.9267

# Building the VGG16 and layering the dense layers and then freezing them

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 4, 4, 512)	14714688
flatten (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 256)	2097408
dense_3 (Dense)	(None, 1)	257

=====  
Total params: 16812353 (64.13 MB)  
Trainable params: 16812353 (64.13 MB)  
Non-trainable params: 0 (0.00 Byte)

In Keras, freezing a network is done by setting its trainable attribute to False:

```
[ ] 1 print('This is the number of trainable weights '
2     'before freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 30

```
[ ] 1 conv_base.trainable = False
```


```
1 print('This is the number of trainable weights '
2     'after freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 4




# Validation Accuracy is at 90.33%

Training and validation accuracy




Training and validation loss



Epoch 30/30

30/30 - 12s - loss: 0.2539 - acc: 0.8933 - val\_loss: 0.2903 - val\_acc: 0.9033 - 12s/  
epoch - 41ms/step

# Fine-Tuning



**Conv block 1:  
frozen**

**Conv block 2:  
frozen**

**Conv block 3:  
frozen**

**Conv block 4:  
frozen**

**We fine-tune  
Conv block 5**

**We fine-tune  
our own  
fully-connected  
classifier**

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14714688 (56.13 MB)		
Trainable params: 0 (0.00 Byte)		
Non-trainable params: 14714688 (56.13 MB)		

# Fine-tune the last 3 convolutional layers

Let's set this up, starting from where we left off in the previous example:

```
[ ] 1 conv_base.trainable = True
    2
    3 #set_trainable = False
    4 for layer in conv_base.layers:
    5     set_trainable = False
    6     if layer.name == 'block5_conv1' or layer.name == 'block5_conv2' or layer.name == 'block5_conv3':
    7         set_trainable = True
    8     if set_trainable:
    9         layer.trainable = True
    10    else:
    11        layer.trainable = False
```

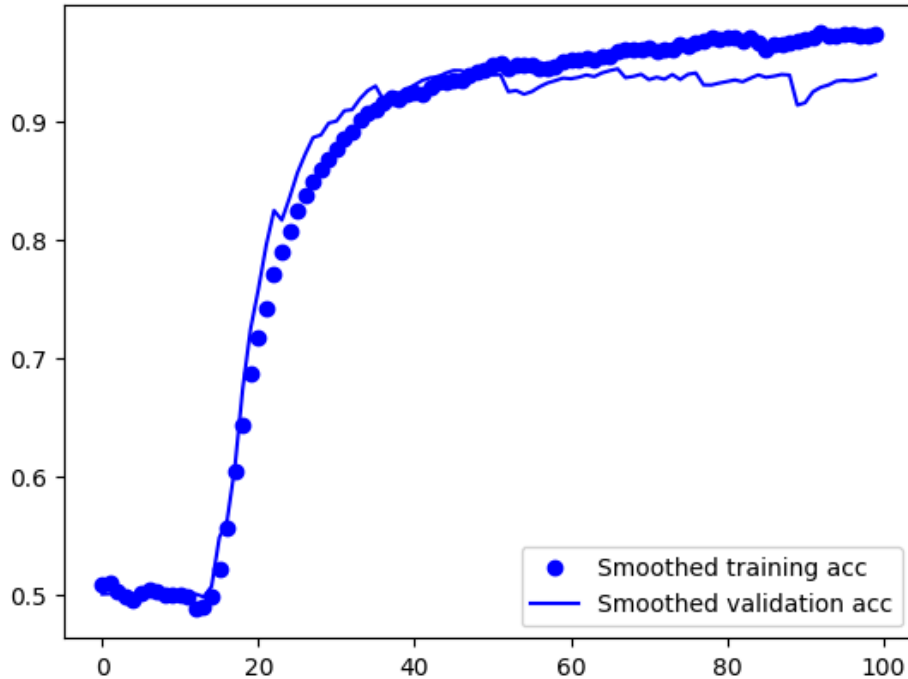
```
[ ] 1 for layer in conv_base.layers:
    2     print(layer.name, ': ', layer.trainable)
    3
```

```
input_1 : False
block1_conv1 : False
block1_conv2 : False
block1_pool : False
block2_conv1 : False
block2_conv2 : False
block2_pool : False
block3_conv1 : False
block3_conv2 : False
block3_conv3 : False
block3_pool : False
block4_conv1 : False
block4_conv2 : False
block4_conv3 : False
block4_pool : False
block5_conv1 : True
block5_conv2 : True
block5_conv3 : True
```

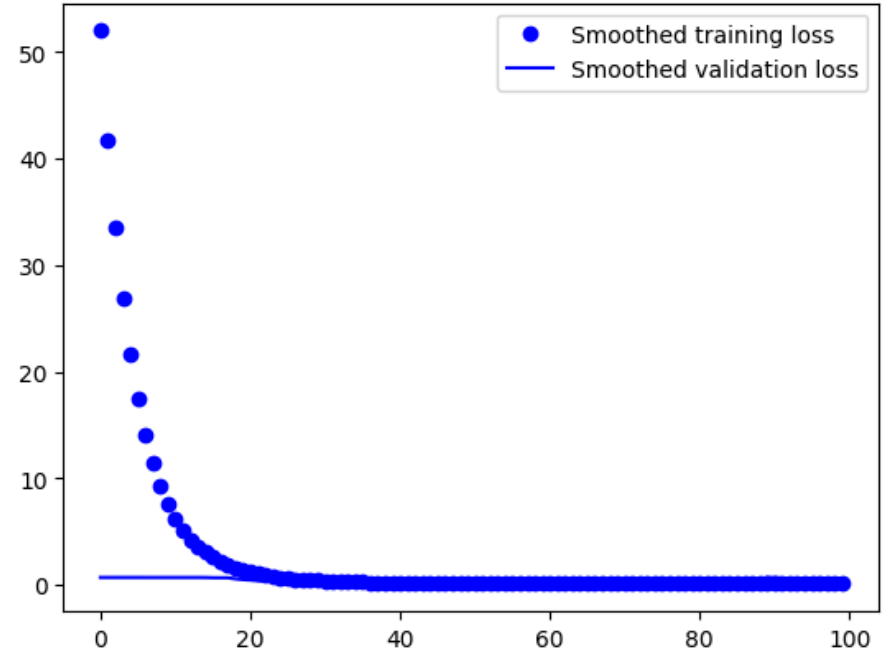
Screenshot

# Test Accuracy 94.04% & Validation Accuracy 95%

Training and validation accuracy



Training and validation loss



Epoch 100/100

30/30 [=====] - 6s 217ms/step - loss: 0.0695 - acc: 0.9767  
- val\_loss: 0.5351 - val\_acc: 0.9500

```
[ ] 1 test_generator = test_datagen.flow_from_directory(  
2     test_dir,  
3     target_size=(150, 150),  
4     batch_size=20,  
5     class_mode='binary')  
6  
7 test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)  
8 print('test acc:', test_acc)
```

Found 806 images belonging to 2 classes.

<ipython-input-27-19f8443b6c42>:7: UserWarning: `Model.evaluate\_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.

test\_loss, test\_acc = model.evaluate\_generator(test\_generator, steps=50)

WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps\_per\_epoch \* epochs` batches (in this case, 50 batches).  
test acc: 0.940446674823761

## Part 3 - Pre-trained Convolutional Neural Networks - Xception

# Pretrained Convolutional Neural Networks - Xception

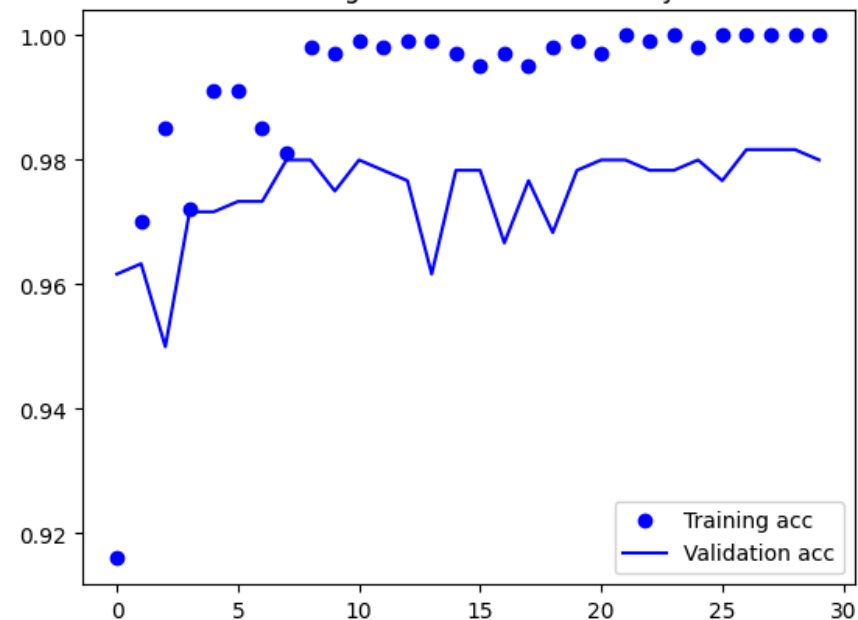
Model: "xception"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 299, 299, 3)]	0	[]
block1_conv1 (Conv2D)	(None, 149, 149, 32)	864	['input_1[0][0]']
block1_conv1_bn (BatchNormalization)	(None, 149, 149, 32)	128	['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, 149, 149, 32)	0	['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, 147, 147, 64)	18432	['block1_conv1_act[0][0]']
block1_conv2_bn (BatchNormalization)	(None, 147, 147, 64)	256	['block1_conv2[0][0]']
block1_conv2_act (Activation)	(None, 147, 147, 64)	0	['block1_conv2_bn[0][0]']
block2_sepconv1 (Separable Conv2D)	(None, 147, 147, 128)	8768	['block1_conv2_act[0][0]']
block2_sepconv1_bn (BatchNormalization)	(None, 147, 147, 128)	512	['block2_sepconv1[0][0]']
block2_sepconv2_act (Activation)	(None, 147, 147, 128)	0	['block2_sepconv1_bn[0][0]']
.....			
block14_sepconv1_act (Activation)	(None, 10, 10, 1536)	0	['block14_sepconv1_bn[0][0]']
block14_sepconv2 (Separable Conv2D)	(None, 10, 10, 2048)	3159552	['block14_sepconv1_act[0][0]']
block14_sepconv2_bn (Batch Normalization)	(None, 10, 10, 2048)	8192	['block14_sepconv2[0][0]']
block14_sepconv2_act (Activation)	(None, 10, 10, 2048)	0	['block14_sepconv2_bn[0][0]']

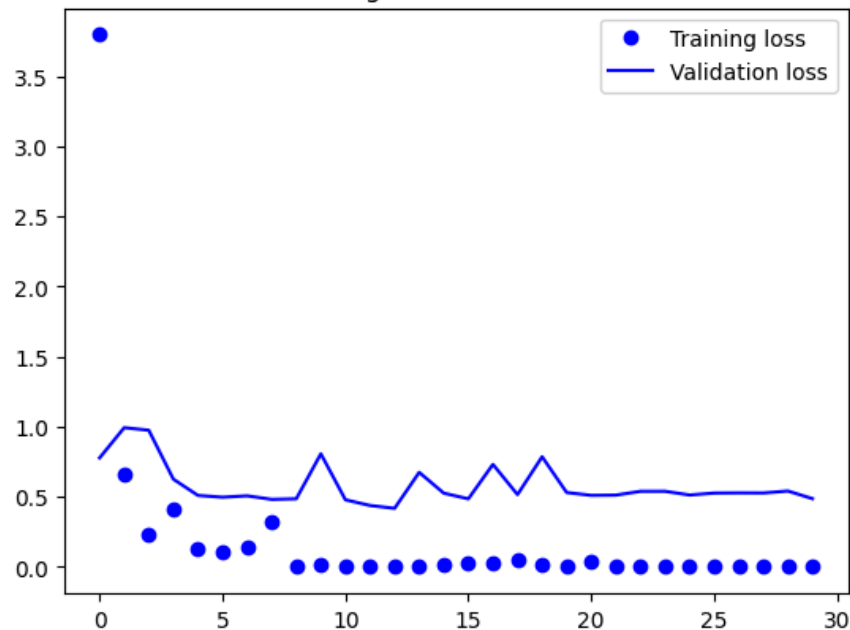
=====  
Total params: 20861480 (79.58 MB)  
Trainable params: 20806952 (79.37 MB)  
Non-trainable params: 54528 (213.00 KB)

# Validation Accuracy 98%

Training and validation accuracy



Training and validation loss



Epoch 30/30

50/50 [=====] - 1s 19ms/step - loss: 1.7528e-04 - acc: 1.0000 - val\_loss: 0.4869 - val\_acc: 0.9800

# Building the Exception and layering the dense layers and then freezing them

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
exception (Functional)	(None, 10, 10, 2048)	20861480
flatten (Flatten)	(None, 204800)	0
dense_2 (Dense)	(None, 256)	52429056
dense_3 (Dense)	(None, 1)	257

Total params: 73290793 (279.58 MB)

Trainable params: 73236265 (279.37 MB)

Non-trainable params: 54528 (213.00 KB)

```
[ ] 1 print('This is the number of trainable weights '
2     '| | | 'before freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 158



```
[ ] 1 conv_base.trainable = False
```

```
[ ] 1 print('This is the number of trainable weights '
2     '| | | 'after freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 4



# Validation Accuracy 97.17%



Epoch 30/30  
50/50 - 29s - loss: 0.0867 - acc: 0.9650 - val\_loss: 0.1237 - val\_acc: 0.9717 - 29s/  
epoch - 575ms/step

# Fine Tuning

Model: "xception"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[ (None, 299, 299, 3) ]	0	[ ]
block1_conv1 (Conv2D)	(None, 149, 149, 32)	864	[ 'input_1[0][0]' ]
block1_conv1_bn (BatchNormalization)	(None, 149, 149, 32)	128	[ 'block1_conv1[0][0]' ]
block1_conv1_act (Activation)	(None, 149, 149, 32)	0	[ 'block1_conv1_bn[0][0]' ]
block1_conv2 (Conv2D)	(None, 147, 147, 64)	18432	[ 'block1_conv1_act[0][0]' ]
block1_conv2_bn (BatchNormalization)	(None, 147, 147, 64)	256	[ 'block1_conv2[0][0]' ]
block1_conv2_act (Activation)	(None, 147, 147, 64)	0	[ 'block1_conv2_bn[0][0]' ]
....			
block14_sepconv2 (SeparableConv2D)	(None, 10, 10, 2048)	3159552	[ 'block14_sepconv1_act[0][0]' ]
block14_sepconv2_bn (BatchNormalization)	(None, 10, 10, 2048)	8192	[ 'block14_sepconv2[0][0]' ]
block14_sepconv2_act (Activation)	(None, 10, 10, 2048)	0	[ 'block14_sepconv2_bn[0][0]' ]
=====			
Total params: 20861480 (79.58 MB)			
Trainable params: 0 (0.00 Byte)			
Non-trainable params: 20861480 (79.58 MB)			

# Fine Tune the last 2 Separable Conv. Layers

Fine Tune the last two Separable Convolutional layers in the last block of the Xception Model.

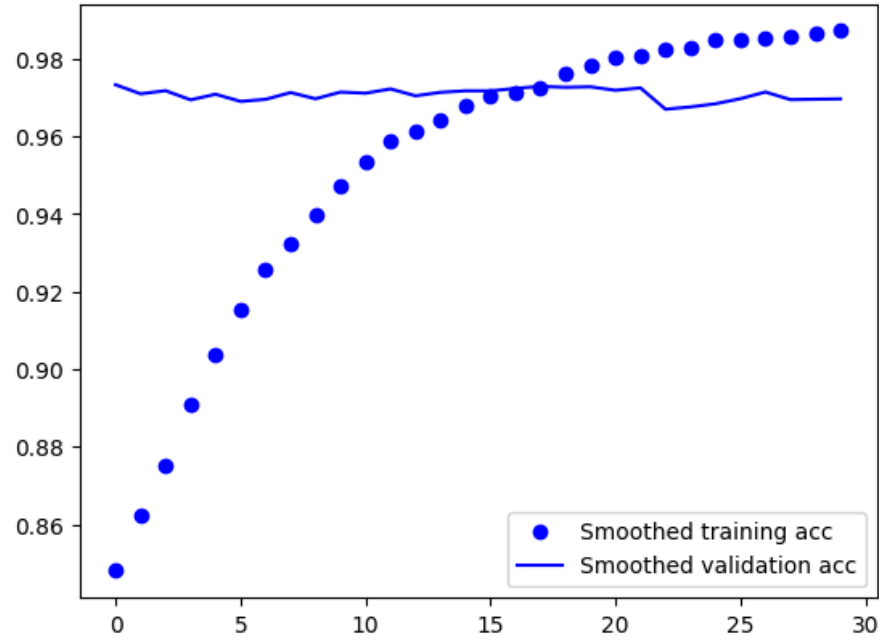
```
[ ] 1 conv_base.trainable = True
    2
    3 #set_trainable = False
    4 for layer in conv_base.layers:
    5     set_trainable = False
    6     if layer.name == 'block14_sepconv1' or layer.name == 'block14_sepconv2' :
    7         set_trainable = True
    8     if set_trainable:
    9         layer.trainable = True
   10     else:
   11         layer.trainable = False
```

```
[ ] 1 for layer in conv_base.layers:
    2     print(layer.name, ': ', layer.trainable)
```

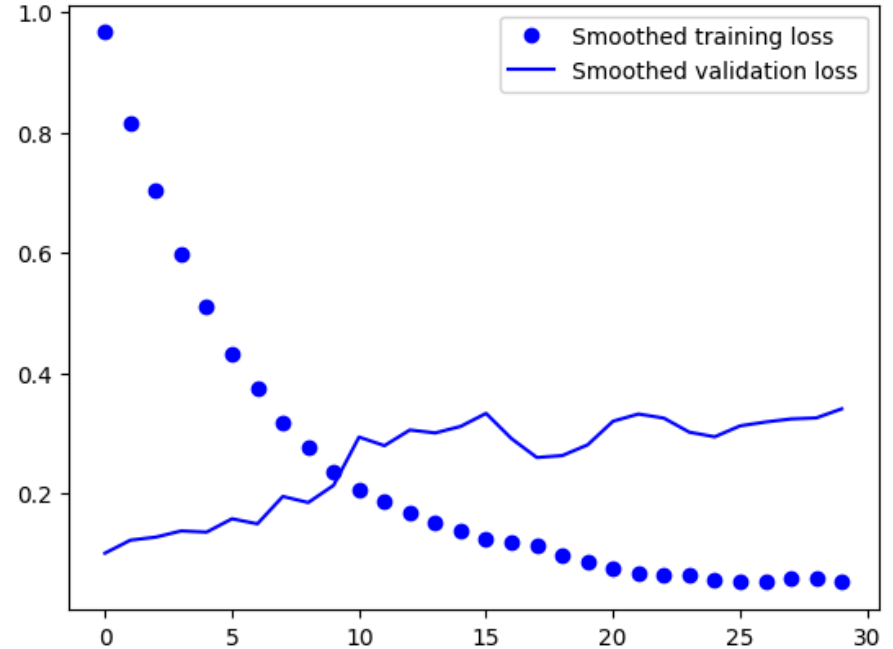
```
-----
add_10 : False
block13_sepconv1_act : False
block13_sepconv1 : False
block13_sepconv1_bn : False
block13_sepconv2_act : False
block13_sepconv2 : False
block13_sepconv2_bn : False
conv2d_3 : False
block13_pool : False
batch_normalization_3 : False
add_11 : False
block14_sepconv1 : True
block14_sepconv1_bn : False
block14_sepconv1_act : False
block14_sepconv2 : True
block14_sepconv2_bn : False
block14_sepconv2_act : False
```

# Accuracy is extremely high with Xception model and fine tuning it at 98.13% vs that of VGG16 at 94.04%

Training and validation accuracy



Training and validation loss



Epoch 30/30

50/50 [=====] - 31s 628ms/step - loss: 0.0295 - acc: 0.9900  
- val\_loss: 0.4008 - val\_acc: 0.9700

```
1 test_generator = test_datagen.flow_from_directory(  
2     test_dir,  
3     target_size=(299, 299),  
4     batch_size=20,  
5     class_mode='binary')  
6  
7 test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)  
8 print('test acc:', test_acc)
```

Found 806 images belonging to 2 classes.

<ipython-input-24-1a4d34cbe021>:7: UserWarning: `Model.evaluate\_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.  
test\_loss, test\_acc = model.evaluate\_generator(test\_generator, steps=50)

WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps\_per\_epoch \* epochs` batches (in this case, 50 batches).  
test acc: 0.981389582157135

# YouTube URLs, Last Page

- Two minute (short): <https://youtu.be/uqDYF-L85D8>
- 15 minutes (long): <https://youtu.be/GPKKnFP IraM>